



tail -f

Fabian Stäber
Munich Gophers Meetup
13 October 2016

Copy-and-Paste from Go Tutorial

```
func main() {  
    file, err := os.Open("file.log")  
    if err != nil {  
        log.Fatal(err)  
    }  
    defer file.Close()  
    reader := bufio.NewReader(file)  
    for {  
        line, err := reader.ReadString('\n')  
        if err != nil {  
            break  
        }  
        fmt.Print(string(line))  
    }  
}
```

Copy-and-Paste from Go Tutorial

```
func main() {  
    file, err := os.Open("file.log")  
    if err != nil {  
        log.Fatal(err)  
    }  
    defer file.Close()  
    reader := bufio.NewReader(file)  
    for {  
        line, err := reader.ReadString('\n')  
        if err != nil {  
            break  
        }  
        fmt.Print(string(line))  
    }  
}
```

This implements 'cat'.

How to extend this to 'tail -f'?

Copy-and-Paste from Go Tutorial

```
func main() {  
    // ...  
    for {  
        line, err := reader.ReadString('\n')  
        if err != nil {  
            break  
        }  
        fmt.Print(string(line))  
    }  
}
```

My first 'tail -f'

```
func main() {  
    // ...  
    for {  
        line, err := reader.ReadString('\n')  
        if err != nil {  
            if err == io.EOF {  
                time.Sleep(1 * time.Second)  
            } else {  
                break  
            }  
        }  
        fmt.Print(string(line))  
    }  
}
```

Logrotate

Demo: Execute these commands while the tailer is running.

```
mv file.log file.log.1  
echo example log line >> file.log  
echo example log line >> file.log.1  
rm file.log.1
```

Note: Cannot do this on Windows, because the file is “in use by another process”.

Logrotate Configurations

Move the old file and create a new one.

```
mv logfile logfile.1  
:> logfile  
echo 'next log line' >> logfile
```

Copy the old file and truncate the original copy.

```
cp logfile logfile.1  
:> logfile  
echo 'next log line' >> logfile
```

How to deal with logrotate

```
if err == io.EOF {  
    time.Sleep(1 * time.Second)  
    // check if file was truncated  
    // check if file was moved  
    // close and re-open if necessary  
}
```



Need to do this.

How to deal with logrotate

```
if err == io.EOF {  
    time.Sleep(1 * time.Second)  
    // check if file was truncated  
    // check if file was moved  
    // close and re-open if necessary  
}
```

Need to do this.



How to do it.



```
fileInfo, err = file.Stat()  
if fileInfo.Size() < curReadPos {  
    // truncated  
}  
if ! os.SameFile(fileInfo, curFileInfo) {  
    // moved  
}
```

FileBeat: github.com/elastic/beats

Pro

- Reasonable and simple

Con

- Requires weird trade-offs in configuration

`backoff, backoff_factor, max_backoff`
`close_eof, close_inactive, close_old, close_removed, close_renamed`

- There must be a way to do this without polling

fsnotify: github.com/fsnotify/fsnotify

Abstract File System notifications.

Used in:

- github.com/hpcloud/tail
- github.com/google/mtail

fsnotify: github.com/fsnotify/fsnotify

```
watcher, err := fsnotify.NewWatcher()
go func() {
    for {
        select {
        case event := <-watcher.Events:
            // ...
        case err := <-watcher.Errors:
            // ...
        }
    }
}
watcher.Add("file.log")
```

Events:

- Create
- Write
- Remove
- Chmod

fsnotify: github.com/fsnotify/fsnotify

Linux inotify() system call mapping:

unix.IN_CREATE, unix.IN_MOVED_TO	fsnotify.Create
unix.IN_DELETE_SELF, unix.IN_DELETE	fsnotify.Remove
unix.IN_MODIFY	fsnotify.Write
unix.IN_MOVE_SELF, unix.IN_MOVED_FROM	fsnotify.Rename
unix.IN_ATTRIB	fsnotify.Chmod
others	ignored

truncate



Fsnotify watches the **directory** where the log file is located.

fsnotify: github.com/fsnotify/fsnotify

BSD kevent() system call mapping:

???	fsnotify.Create
unix.NOTE_DELETE	fsnotify.Remove
unix.NOTE_WRITE	fsnotify.Write
unix.NOTE_RENAME	fsnotify.Rename
unix.NOTE_ATTRIB	fsnotify.Chmod
others	ignored

truncate



Fsnotify “simulates” recursive directory watches.

fsnotify: github.com/fsnotify/fsnotify

OS-specific code needed:

- Make up for missing WRITE events due to races on BSD
- Figure out if file was truncated on fsnotify.Write and fsnotify.Chmod
- Must close file on Windows, can only watch open files on BSD?
- ...

OS-specific corner cases must be found with trial-and-error.

Rigid testing needed, use Travis CI for OS X and Linux, AppVeyor for Windows.

Most fsnotify-based tools ignore this and focus on Linux.

fsnotify: github.com/fsnotify/fsnotify

OS specific
file system
events



OS independent
fsnotify events



OS specific
interpretation of
fsnotify events

Better use system calls directly.

BSD and Linux example

BSD: kevent()

```
fd = syscall.Kqueue()
```

```
go func() {  
    for {  
        syscall.Kevent(fd, ...)   
        // ...  
        logData <- data  
    }  
}
```

Linux: inotify()

```
fd = syscall.InotifyInit1(...)   
syscall.InotifyAddWatch(fd, ...)
```

```
go func() {  
    for {  
        syscall.Read(fd, ...)   
        // ...  
        logData <- data  
    }  
}
```

How to shut down

Producer: Event loop

```
go func() {  
    for {  
        err = syscall.Read(fd, ...)  
        if err == interrupted {  
            return  
        }  
        // ...  
        logData <- data  
    }  
}
```

Consumer: Log line processor

```
select {  
    case data := <- logData:  
        // do something with data  
    case <- quit:  
        // interrupt the system call  
}
```

How to interrupt read() ?

How to shut down

BSD: kevent()

`syscall.Close(fd)`

Linux: inotify()

`syscall.InotifyRmWatch(fd, ...)`

How to shut down

Producer: Event loop

```
go func() {  
    for {  
        err = syscall.Read(fd, ...)  
        if err == interrupted {  
            return  
        }  
        // ...  
        logData <- data  
    }  
}
```

Consumer: Log line processor

```
select {  
    case data := <- logData:  
        // do something with data  
    case <- quit:  
        // interrupt the system call  
}
```

How to shut down

Producer: Event loop

```
go func() {  
    for {  
        err = syscall.Read(fd, ...)  
        if err == interrupted {  
            return  
        }  
        // ...  
        logData <- data  
    }  
}
```

Consumer: Log line processor

```
select {  
    case data := <- logData:  
        // do something with data  
    case <- quit:  
        // interrupt the system call  
}
```

What if the producer hangs here
when read() is interrupted?

How to shut down

Producer: Event loop

```
go func() {  
    for {  
        // ...  
        logData <- data  
    }  
}
```

Consumer: Log line processor

```
select {  
    case data := <- logData:  
        // do something with data  
    case <- quit:  
        // interrupt the system call  
}
```

How to shut down

Producer: Event loop

```
go func() {  
    for {  
        // ...  
        select {  
            case logData <- data:  
            case <- done:  
                return  
        }  
    }  
}
```

Consumer: Log line processor

```
select {  
    case data := <- logData:  
        // do something with data  
    case <- quit:  
        close(done)  
        // interrupt the system call  
}
```

Thank you!



github.com/fstab/grok_exporter
[package tailer/](#)



[@fstabr](https://twitter.com/fstabr)



<http://www.consol.de>